

m7vxdqs4n

April 18, 2025

0.1 Linear regression by using Deep Neural network: Implement Boston housing price prediction problem by Linear regression using Deep Neural Network. Use Boston House Price prediction Dataset.

```
[1]: import tensorflow as tf
      from tensorflow import keras
      from tensorflow.keras import layers
      import numpy as np
      import pandas as pd
      import matplotlib.pyplot as plt
      from sklearn.preprocessing import StandardScaler
      from sklearn.model_selection import train_test_split
      from tensorflow.keras.optimizers import Adam
```

```
[2]: df = pd.read_csv("HousingData.csv")
      df.head()
```

```
[2]:
```

| | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO | \ |
|---|---------|------|-------|------|-------|-------|------|--------|-----|-----|---------|---|
| 0 | 0.00632 | 18.0 | 2.31 | 0.0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1 | 296 | 15.3 | |
| 1 | 0.02731 | 0.0 | 7.07 | 0.0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2 | 242 | 17.8 | |
| 2 | 0.02729 | 0.0 | 7.07 | 0.0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2 | 242 | 17.8 | |
| 3 | 0.03237 | 0.0 | 2.18 | 0.0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3 | 222 | 18.7 | |
| 4 | 0.06905 | 0.0 | 2.18 | 0.0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3 | 222 | 18.7 | |

| | B | LSTAT | MEDV |
|---|--------|-------|------|
| 0 | 396.90 | 4.98 | 24.0 |
| 1 | 396.90 | 9.14 | 21.6 |
| 2 | 392.83 | 4.03 | 34.7 |
| 3 | 394.63 | 2.94 | 33.4 |
| 4 | 396.90 | NaN | 36.2 |

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---
```

```

0   CRIM      486 non-null   float64
1   ZN        486 non-null   float64
2   INDUS     486 non-null   float64
3   CHAS      486 non-null   float64
4   NOX       506 non-null   float64
5   RM        506 non-null   float64
6   AGE       486 non-null   float64
7   DIS       506 non-null   float64
8   RAD       506 non-null   int64
9   TAX       506 non-null   int64
10  PTRATIO   506 non-null   float64
11  B         506 non-null   float64
12  LSTAT     486 non-null   float64
13  MEDV      506 non-null   float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB

```

```
[4]: df.fillna(df.mean(), inplace=True)
```

```
[5]: X = df.drop(columns=['MEDV'])
     y = df['MEDV']
```

```
[6]: scaler = StandardScaler()
     X = scaler.fit_transform(X)
```

```
[7]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
     ↪ random_state=42)
```

```
[8]: model = keras.Sequential([
     keras.Input(shape=(X.shape[1],)), # Explicit Input Layer
     layers.Dense(64, activation='relu'), # Hidden Layer 1
     layers.Dense(32, activation='relu'), # Hidden Layer 2
     layers.Dense(1, activation='linear') # Output layer (Regression)
])
```

```
[9]: model.compile(optimizer=Adam(learning_rate=0.01), loss='mse', metrics=['mae'])
```

```
[10]: history = model.fit(X_train, y_train, epochs=100, validation_data=(X_test,
     ↪ y_test), batch_size=16, verbose=1)
```

```

Epoch 1/100
26/26          1s 8ms/step - loss:
440.8634 - mae: 18.5338 - val_loss: 57.2412 - val_mae: 5.5565
Epoch 2/100
26/26          0s 2ms/step - loss:
45.3372 - mae: 5.1544 - val_loss: 21.6366 - val_mae: 3.2939
Epoch 3/100
26/26          0s 3ms/step - loss:

```

20.0994 - mae: 3.3215 - val_loss: 17.6297 - val_mae: 2.8288
 Epoch 4/100
 26/26 0s 4ms/step - loss:
 17.6236 - mae: 3.0781 - val_loss: 14.9943 - val_mae: 2.6249
 Epoch 5/100
 26/26 0s 3ms/step - loss:
 14.2821 - mae: 2.7823 - val_loss: 14.5623 - val_mae: 2.3381
 Epoch 6/100
 26/26 0s 2ms/step - loss:
 13.6029 - mae: 2.6609 - val_loss: 12.9750 - val_mae: 2.4203
 Epoch 7/100
 26/26 0s 3ms/step - loss:
 12.8029 - mae: 2.5281 - val_loss: 13.4118 - val_mae: 2.3723
 Epoch 8/100
 26/26 0s 3ms/step - loss:
 14.5553 - mae: 2.7077 - val_loss: 13.8585 - val_mae: 2.4914
 Epoch 9/100
 26/26 0s 3ms/step - loss:
 13.7882 - mae: 2.5705 - val_loss: 15.7087 - val_mae: 2.6962
 Epoch 10/100
 26/26 0s 3ms/step - loss:
 10.6948 - mae: 2.4719 - val_loss: 15.3429 - val_mae: 2.6446
 Epoch 11/100
 26/26 0s 3ms/step - loss:
 14.5759 - mae: 2.7952 - val_loss: 14.6829 - val_mae: 2.7426
 Epoch 12/100
 26/26 0s 3ms/step - loss:
 12.2882 - mae: 2.5158 - val_loss: 12.0716 - val_mae: 2.3263
 Epoch 13/100
 26/26 0s 3ms/step - loss:
 9.0725 - mae: 2.1205 - val_loss: 14.1640 - val_mae: 2.4194
 Epoch 14/100
 26/26 0s 3ms/step - loss:
 8.8355 - mae: 2.2726 - val_loss: 14.3075 - val_mae: 2.5564
 Epoch 15/100
 26/26 0s 3ms/step - loss:
 10.8129 - mae: 2.4245 - val_loss: 13.6643 - val_mae: 2.4626
 Epoch 16/100
 26/26 0s 4ms/step - loss:
 9.6793 - mae: 2.4046 - val_loss: 13.4059 - val_mae: 2.5891
 Epoch 17/100
 26/26 0s 2ms/step - loss:
 10.1957 - mae: 2.4241 - val_loss: 15.5057 - val_mae: 2.7194
 Epoch 18/100
 26/26 0s 2ms/step - loss:
 10.0053 - mae: 2.4237 - val_loss: 14.4027 - val_mae: 2.6199
 Epoch 19/100
 26/26 0s 3ms/step - loss:

8.8556 - mae: 2.3036 - val_loss: 15.0273 - val_mae: 2.5188
 Epoch 20/100
 26/26 0s 2ms/step - loss:
 11.1176 - mae: 2.5842 - val_loss: 13.6674 - val_mae: 2.4945
 Epoch 21/100
 26/26 0s 3ms/step - loss:
 8.4296 - mae: 2.2924 - val_loss: 12.5635 - val_mae: 2.4181
 Epoch 22/100
 26/26 0s 2ms/step - loss:
 8.5037 - mae: 2.1967 - val_loss: 13.1299 - val_mae: 2.3947
 Epoch 23/100
 26/26 0s 2ms/step - loss:
 9.4739 - mae: 2.2657 - val_loss: 12.7473 - val_mae: 2.5449
 Epoch 24/100
 26/26 0s 2ms/step - loss:
 8.6330 - mae: 2.2670 - val_loss: 13.1207 - val_mae: 2.4739
 Epoch 25/100
 26/26 0s 3ms/step - loss:
 8.5585 - mae: 2.2271 - val_loss: 11.9104 - val_mae: 2.2847
 Epoch 26/100
 26/26 0s 2ms/step - loss:
 8.1119 - mae: 2.1407 - val_loss: 13.2038 - val_mae: 2.5670
 Epoch 27/100
 26/26 0s 2ms/step - loss:
 8.2387 - mae: 2.2445 - val_loss: 12.3271 - val_mae: 2.4717
 Epoch 28/100
 26/26 0s 2ms/step - loss:
 10.2736 - mae: 2.4373 - val_loss: 13.2581 - val_mae: 2.4999
 Epoch 29/100
 26/26 0s 3ms/step - loss:
 9.6498 - mae: 2.2557 - val_loss: 13.1113 - val_mae: 2.5810
 Epoch 30/100
 26/26 0s 2ms/step - loss:
 8.0430 - mae: 2.0943 - val_loss: 16.5543 - val_mae: 2.6920
 Epoch 31/100
 26/26 0s 2ms/step - loss:
 9.1912 - mae: 2.3014 - val_loss: 13.0773 - val_mae: 2.4636
 Epoch 32/100
 26/26 0s 2ms/step - loss:
 6.8917 - mae: 2.0088 - val_loss: 13.4819 - val_mae: 2.4822
 Epoch 33/100
 26/26 0s 2ms/step - loss:
 8.8703 - mae: 2.3257 - val_loss: 13.6802 - val_mae: 2.7276
 Epoch 34/100
 26/26 0s 2ms/step - loss:
 9.3434 - mae: 2.3046 - val_loss: 11.9395 - val_mae: 2.2687
 Epoch 35/100
 26/26 0s 3ms/step - loss:

6.3509 - mae: 1.9543 - val_loss: 13.7159 - val_mae: 2.5541
 Epoch 36/100
 26/26 0s 2ms/step - loss:
 8.4139 - mae: 2.1992 - val_loss: 14.7841 - val_mae: 2.5076
 Epoch 37/100
 26/26 0s 2ms/step - loss:
 5.0453 - mae: 1.6810 - val_loss: 11.9874 - val_mae: 2.4051
 Epoch 38/100
 26/26 0s 2ms/step - loss:
 7.4600 - mae: 2.0501 - val_loss: 13.0326 - val_mae: 2.4529
 Epoch 39/100
 26/26 0s 2ms/step - loss:
 9.1731 - mae: 2.3376 - val_loss: 12.4445 - val_mae: 2.4682
 Epoch 40/100
 26/26 0s 2ms/step - loss:
 9.4816 - mae: 2.3027 - val_loss: 11.6693 - val_mae: 2.4574
 Epoch 41/100
 26/26 0s 3ms/step - loss:
 7.4163 - mae: 2.0914 - val_loss: 11.5701 - val_mae: 2.3401
 Epoch 42/100
 26/26 0s 1ms/step - loss:
 6.8931 - mae: 2.0371 - val_loss: 13.6548 - val_mae: 2.4897
 Epoch 43/100
 26/26 0s 2ms/step - loss:
 7.0245 - mae: 2.0263 - val_loss: 12.0620 - val_mae: 2.4173
 Epoch 44/100
 26/26 0s 2ms/step - loss:
 6.6387 - mae: 1.9324 - val_loss: 12.6175 - val_mae: 2.3677
 Epoch 45/100
 26/26 0s 2ms/step - loss:
 4.9877 - mae: 1.6471 - val_loss: 11.8034 - val_mae: 2.3359
 Epoch 46/100
 26/26 0s 2ms/step - loss:
 6.4296 - mae: 1.8254 - val_loss: 14.0951 - val_mae: 2.4861
 Epoch 47/100
 26/26 0s 2ms/step - loss:
 6.7578 - mae: 1.8863 - val_loss: 12.3211 - val_mae: 2.4534
 Epoch 48/100
 26/26 0s 2ms/step - loss:
 7.8392 - mae: 2.1540 - val_loss: 19.0252 - val_mae: 2.9707
 Epoch 49/100
 26/26 0s 2ms/step - loss:
 7.9342 - mae: 2.1704 - val_loss: 11.3792 - val_mae: 2.3989
 Epoch 50/100
 26/26 0s 3ms/step - loss:
 8.6149 - mae: 2.0973 - val_loss: 12.0656 - val_mae: 2.4236
 Epoch 51/100
 26/26 0s 2ms/step - loss:

6.7459 - mae: 1.9316 - val_loss: 12.6347 - val_mae: 2.4684
 Epoch 52/100
 26/26 0s 2ms/step - loss:
 8.0481 - mae: 2.1807 - val_loss: 12.8758 - val_mae: 2.3664
 Epoch 53/100
 26/26 0s 1ms/step - loss:
 6.2552 - mae: 1.8587 - val_loss: 12.6247 - val_mae: 2.7060
 Epoch 54/100
 26/26 0s 2ms/step - loss:
 6.8880 - mae: 2.0309 - val_loss: 11.8173 - val_mae: 2.3182
 Epoch 55/100
 26/26 0s 2ms/step - loss:
 5.4851 - mae: 1.8217 - val_loss: 11.3842 - val_mae: 2.2862
 Epoch 56/100
 26/26 0s 2ms/step - loss:
 5.5749 - mae: 1.8342 - val_loss: 13.7648 - val_mae: 2.4022
 Epoch 57/100
 26/26 0s 2ms/step - loss:
 7.1245 - mae: 1.9695 - val_loss: 15.4390 - val_mae: 2.6652
 Epoch 58/100
 26/26 0s 2ms/step - loss:
 6.0535 - mae: 1.8371 - val_loss: 12.1460 - val_mae: 2.4247
 Epoch 59/100
 26/26 0s 6ms/step - loss:
 5.2226 - mae: 1.6742 - val_loss: 14.3516 - val_mae: 2.4435
 Epoch 60/100
 26/26 0s 2ms/step - loss:
 5.8027 - mae: 1.7942 - val_loss: 11.3104 - val_mae: 2.3171
 Epoch 61/100
 26/26 0s 3ms/step - loss:
 5.5879 - mae: 1.7765 - val_loss: 13.8340 - val_mae: 2.4806
 Epoch 62/100
 26/26 0s 3ms/step - loss:
 5.5349 - mae: 1.7566 - val_loss: 13.4828 - val_mae: 2.4362
 Epoch 63/100
 26/26 0s 2ms/step - loss:
 6.2231 - mae: 1.8632 - val_loss: 16.7503 - val_mae: 2.9819
 Epoch 64/100
 26/26 0s 2ms/step - loss:
 7.2343 - mae: 2.0410 - val_loss: 12.8258 - val_mae: 2.7086
 Epoch 65/100
 26/26 0s 2ms/step - loss:
 6.7012 - mae: 1.9879 - val_loss: 12.6499 - val_mae: 2.4648
 Epoch 66/100
 26/26 0s 2ms/step - loss:
 6.2187 - mae: 1.9058 - val_loss: 16.2132 - val_mae: 2.6801
 Epoch 67/100
 26/26 0s 2ms/step - loss:

7.3492 - mae: 2.1787 - val_loss: 12.5516 - val_mae: 2.3207
 Epoch 68/100
 26/26 0s 2ms/step - loss:
 8.4214 - mae: 2.2212 - val_loss: 13.8353 - val_mae: 2.8145
 Epoch 69/100
 26/26 0s 3ms/step - loss:
 7.3707 - mae: 2.1419 - val_loss: 15.5721 - val_mae: 2.8678
 Epoch 70/100
 26/26 0s 2ms/step - loss:
 6.5717 - mae: 2.0063 - val_loss: 13.7526 - val_mae: 2.4975
 Epoch 71/100
 26/26 0s 2ms/step - loss:
 7.7291 - mae: 2.0831 - val_loss: 12.3777 - val_mae: 2.3361
 Epoch 72/100
 26/26 0s 5ms/step - loss:
 5.2945 - mae: 1.7679 - val_loss: 13.3363 - val_mae: 2.7306
 Epoch 73/100
 26/26 0s 2ms/step - loss:
 5.7156 - mae: 1.8765 - val_loss: 12.8553 - val_mae: 2.5685
 Epoch 74/100
 26/26 0s 3ms/step - loss:
 5.7010 - mae: 1.8624 - val_loss: 14.7535 - val_mae: 2.3943
 Epoch 75/100
 26/26 0s 2ms/step - loss:
 5.3147 - mae: 1.7076 - val_loss: 13.4687 - val_mae: 2.3820
 Epoch 76/100
 26/26 0s 2ms/step - loss:
 4.5659 - mae: 1.5840 - val_loss: 13.4049 - val_mae: 2.4064
 Epoch 77/100
 26/26 0s 3ms/step - loss:
 6.5707 - mae: 1.9099 - val_loss: 11.9541 - val_mae: 2.2850
 Epoch 78/100
 26/26 0s 2ms/step - loss:
 5.2888 - mae: 1.7650 - val_loss: 14.1747 - val_mae: 2.6666
 Epoch 79/100
 26/26 0s 2ms/step - loss:
 5.9275 - mae: 1.8753 - val_loss: 12.7242 - val_mae: 2.4817
 Epoch 80/100
 26/26 0s 2ms/step - loss:
 4.5076 - mae: 1.6957 - val_loss: 11.2360 - val_mae: 2.3203
 Epoch 81/100
 26/26 0s 2ms/step - loss:
 5.9954 - mae: 1.8458 - val_loss: 14.4281 - val_mae: 2.4375
 Epoch 82/100
 26/26 0s 2ms/step - loss:
 5.1949 - mae: 1.7030 - val_loss: 11.7536 - val_mae: 2.2650
 Epoch 83/100
 26/26 0s 1ms/step - loss:

5.1825 - mae: 1.7165 - val_loss: 12.7035 - val_mae: 2.3677
 Epoch 84/100
 26/26 0s 3ms/step - loss:
 5.5240 - mae: 1.7834 - val_loss: 12.4507 - val_mae: 2.4028
 Epoch 85/100
 26/26 0s 3ms/step - loss:
 5.3263 - mae: 1.6635 - val_loss: 11.4906 - val_mae: 2.4123
 Epoch 86/100
 26/26 0s 2ms/step - loss:
 5.4103 - mae: 1.7492 - val_loss: 11.7925 - val_mae: 2.2501
 Epoch 87/100
 26/26 0s 2ms/step - loss:
 5.1477 - mae: 1.6802 - val_loss: 12.8947 - val_mae: 2.3426
 Epoch 88/100
 26/26 0s 2ms/step - loss:
 4.7069 - mae: 1.6337 - val_loss: 12.2875 - val_mae: 2.2313
 Epoch 89/100
 26/26 0s 2ms/step - loss:
 6.4071 - mae: 1.8286 - val_loss: 11.5631 - val_mae: 2.5627
 Epoch 90/100
 26/26 0s 3ms/step - loss:
 5.6254 - mae: 1.8210 - val_loss: 12.5988 - val_mae: 2.2553
 Epoch 91/100
 26/26 0s 1ms/step - loss:
 4.3570 - mae: 1.5867 - val_loss: 13.4870 - val_mae: 2.4412
 Epoch 92/100
 26/26 0s 2ms/step - loss:
 3.7779 - mae: 1.4996 - val_loss: 13.0153 - val_mae: 2.3395
 Epoch 93/100
 26/26 0s 2ms/step - loss:
 5.1520 - mae: 1.7093 - val_loss: 13.6739 - val_mae: 2.5711
 Epoch 94/100
 26/26 0s 2ms/step - loss:
 5.6690 - mae: 1.7843 - val_loss: 12.1259 - val_mae: 2.3143
 Epoch 95/100
 26/26 0s 3ms/step - loss:
 5.1884 - mae: 1.7702 - val_loss: 12.9545 - val_mae: 2.4932
 Epoch 96/100
 26/26 0s 2ms/step - loss:
 4.8350 - mae: 1.6868 - val_loss: 11.5383 - val_mae: 2.2410
 Epoch 97/100
 26/26 0s 2ms/step - loss:
 4.8329 - mae: 1.6504 - val_loss: 14.1497 - val_mae: 2.6293
 Epoch 98/100
 26/26 0s 2ms/step - loss:
 5.1155 - mae: 1.7440 - val_loss: 11.5848 - val_mae: 2.3331
 Epoch 99/100
 26/26 0s 2ms/step - loss:


```
5.3956 - mae: 1.7245 - val_loss: 11.9814 - val_mae: 2.3072
Epoch 100/100
26/26          0s 2ms/step - loss:
3.7539 - mae: 1.5173 - val_loss: 11.3920 - val_mae: 2.2617
```


```
[11]: loss, mae = model.evaluate(X_test, y_test)
      print(f"Test Loss (MSE): {loss}")
      print(f"Test Mean Absolute Error (MAE): {mae}")
```

```
4/4          0s 0s/step - loss:
8.9360 - mae: 2.1187
Test Loss (MSE): 11.391955375671387
Test Mean Absolute Error (MAE): 2.261707305908203
```

```
[12]: predictions = model.predict(X_test)
```

```
4/4          0s 6ms/step
```

```
[13]: plt.scatter(y_test, predictions)
      plt.xlabel("Actual Prices ($1000s)")
      plt.ylabel("Predicted Prices ($1000s)")
      plt.title("Actual vs Predicted House Prices")
      plt.show()
```



wxs98airg

April 18, 2025

0.1 Binary classification using Deep Neural Networks Example: Classify movie reviews into “positive” reviews and “negative” reviews, just based on the text content of the reviews. Use IMDB dataset.

```
[1]: import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import matplotlib.pyplot as plt

[2]: vocab_size = 10000
max_length = 200
(x_train, y_train), (x_test, y_test) = keras.datasets.imdb.
↳load_data(num_words=vocab_size)

[3]: x_train = keras.preprocessing.sequence.pad_sequences(x_train,
↳maxlen=max_length, padding='post')
x_test = keras.preprocessing.sequence.pad_sequences(x_test, maxlen=max_length,
↳padding='post')

[4]: model = keras.Sequential([
    layers.Embedding(input_dim=vocab_size, output_dim=64), # Removed
↳input_length
    layers.Conv1D(32, 5, activation='relu'), # 1D Convolution for feature
↳extraction
    layers.GlobalMaxPooling1D(), # Reduce dimensions
    layers.Dense(64, activation='relu'), # Fully connected layer
    layers.Dense(1, activation='sigmoid') # Output layer (Binary
↳classification)
])

[5]: model.compile(optimizer='adam', loss='binary_crossentropy',
↳metrics=['accuracy'])

[6]: history = model.fit(x_train, y_train, epochs=5, validation_data=(x_test,
↳y_test), batch_size=64, verbose=1)
```

Epoch 1/5

391/391


8s 18ms/step -

```
accuracy: 0.6687 - loss: 0.5754 - val_accuracy: 0.8692 - val_loss: 0.3026
Epoch 2/5
391/391          7s 18ms/step -
accuracy: 0.9185 - loss: 0.2122 - val_accuracy: 0.8845 - val_loss: 0.2767
Epoch 3/5
391/391          7s 18ms/step -
accuracy: 0.9827 - loss: 0.0740 - val_accuracy: 0.8825 - val_loss: 0.3220
Epoch 4/5
391/391          7s 18ms/step -
accuracy: 0.9979 - loss: 0.0180 - val_accuracy: 0.8816 - val_loss: 0.3672
Epoch 5/5
391/391          7s 17ms/step -
accuracy: 1.0000 - loss: 0.0031 - val_accuracy: 0.8836 - val_loss: 0.4037
```

```
[7]: loss, accuracy = model.evaluate(x_test, y_test)
     print(f"Test Accuracy: {accuracy:.4f}")
     print(f"Test Loss: {loss:.4f}")
```

```
782/782          3s 4ms/step -
accuracy: 0.8828 - loss: 0.4049
Test Accuracy: 0.8836
Test Loss: 0.4037
```

```
[8]: plt.plot(history.history['accuracy'], label='Train Accuracy')
     plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
     plt.xlabel('Epochs')
     plt.ylabel('Accuracy')
     plt.legend()
     plt.title('Training and Validation Accuracy')
     plt.show()
```



rbezg8v4z

April 18, 2025

0.1 Convolutional neural network (CNN) (Any One from the following)

Use any dataset of plant disease and design a plant disease detection system using CNN.
Use MNIST Fashion Dataset and create a classifier to classify fashion clothing into categories.

```
[1]: import pandas as pd
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers, models
```

```
[2]: train_df = pd.read_csv('fashion-mnist_train.csv')
test_df = pd.read_csv('fashion-mnist_test.csv')
```

```
[3]: train_df.head(2)
```

```
[3]:   label  pixel1  pixel2  pixel3  pixel4  pixel5  pixel6  pixel7  pixel8  \
0      2        0        0        0        0        0        0        0        0
1      9        0        0        0        0        0        0        0        0

   pixel9  ...  pixel775  pixel776  pixel777  pixel778  pixel779  pixel780  \
0         0  ...         0         0         0         0         0         0
1         0  ...         0         0         0         0         0         0

   pixel781  pixel782  pixel783  pixel784
0          0          0          0          0
1          0          0          0          0

[2 rows x 785 columns]
```

```
[4]: test_df.head(2)
```

```
[4]:   label  pixel1  pixel2  pixel3  pixel4  pixel5  pixel6  pixel7  pixel8  \
0      0        0        0        0        0        0        0        0        9
1      1        0        0        0        0        0        0        0        0

   pixel9  ...  pixel775  pixel776  pixel777  pixel778  pixel779  pixel780  \
0         8  ...      103        87        56          0          0          0
```

```
1      0 ...      34      0      0      0      0      0
```

```
      pixel781 pixel782 pixel783 pixel784
0          0          0          0          0
1          0          0          0          0
```

```
[2 rows x 785 columns]
```

```
[5]: # Split features and labels
x_train = train_df.iloc[:, 1:].values
y_train = train_df.iloc[:, 0].values
```

```
[6]: x_test = test_df.iloc[:, 1:].values
y_test = test_df.iloc[:, 0].values
```

```
[7]: # Normalize pixel values
x_train = x_train / 255.0
x_test = x_test / 255.0
```

```
[8]: # Reshape for CNN: (samples, height, width, channels)
x_train = x_train.reshape(-1, 28, 28, 1)
x_test = x_test.reshape(-1, 28, 28, 1)
```

```
[9]: # Build CNN model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D(2, 2),

    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])
```

C:\Python312\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:107:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
[10]: # Compile model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```
[11]: # Train
model.fit(x_train, y_train, epochs=5, validation_split=0.1)
```

```
Epoch 1/5
1688/1688          13s 7ms/step -
accuracy: 0.7610 - loss: 0.6718 - val_accuracy: 0.8638 - val_loss: 0.3871
Epoch 2/5
1688/1688          12s 7ms/step -
accuracy: 0.8786 - loss: 0.3398 - val_accuracy: 0.8882 - val_loss: 0.3101
Epoch 3/5
1688/1688          13s 7ms/step -
accuracy: 0.8976 - loss: 0.2819 - val_accuracy: 0.8890 - val_loss: 0.3041
Epoch 4/5
1688/1688          13s 8ms/step -
accuracy: 0.9079 - loss: 0.2526 - val_accuracy: 0.8978 - val_loss: 0.2800
Epoch 5/5
1688/1688          12s 7ms/step -
accuracy: 0.9167 - loss: 0.2238 - val_accuracy: 0.8975 - val_loss: 0.2802
```

```
[11]: <keras.src.callbacks.history.History at 0x25f324b7050>
```

```
[12]: # Evaluate
loss, acc = model.evaluate(x_test, y_test)
print(f"\nTest Accuracy: {acc}")
```

```
313/313           1s 4ms/step -
accuracy: 0.9019 - loss: 0.2680
```

```
Test Accuracy: 0.9064000248908997
```

```
[13]: import matplotlib.pyplot as plt
class_names=['T-shirt/
↳top', 'Trouser', 'Pullover', 'Dress', 'Coat', 'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle_
↳boot']
plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(x_train[i], cmap=plt.cm.binary)
    plt.xlabel(class_names[y_train[i]])
plt.show()
```




[]:

//Design and implement Parallel Breadth First Search and Depth First Search based on existing algorithms using OpenMP. Use a Tree or an undirected graph for BFS and DFS .

```
#include <iostream>
#include <vector>
#include <queue>
#include <stack>
#include <omp.h>
using namespace std;

const int N = 6; // Number of nodes
vector<int> graph[N];
bool visited_bfs[N], visited_dfs[N];

// Add edge to undirected graph
void addEdge(int u, int v) {
    graph[u].push_back(v);
    graph[v].push_back(u);
}

// Parallel BFS using OpenMP
void parallelBFS(int start) {
    queue<int> q;
    q.push(start);
    visited_bfs[start] = true;

    while (!q.empty()) {
        int size = q.size();

        #pragma omp parallel for
        for (int i = 0; i < size; i++) {
            int node;
            #pragma omp critical
            {
                node = q.front(); q.pop();
                cout << "BFS visited: " << node << endl;
            }

            for (int neighbor : graph[node]) {
                #pragma omp critical
                {
                    if (!visited_bfs[neighbor]) {
                        visited_bfs[neighbor] = true;
                        q.push(neighbor);
                    }
                }
            }
        }
    }
}

// Parallel DFS using OpenMP
void parallelDFS(int start) {
```

```

stack<int> s;
s.push(start);
visited_dfs[start] = true;

while (!s.empty()) {
    int node;
    #pragma omp critical
    {
        node = s.top(); s.pop();
        cout << "DFS visited: " << node << endl;
    }

    #pragma omp parallel for
    for (int i = 0; i < graph[node].size(); i++) {
        int neighbor = graph[node][i];
        #pragma omp critical
        {
            if (!visited_dfs[neighbor]) {
                visited_dfs[neighbor] = true;
                s.push(neighbor);
            }
        }
    }
}

```

```

int main() {
    addEdge(0, 1);
    addEdge(0, 2);
    addEdge(1, 3);
    addEdge(1, 4);
    addEdge(2, 5);

    cout << "Parallel BFS:\n";
    parallelBFS(0);

    cout << "\nParallel DFS:\n";
    parallelDFS(0);

    return 0;
}

```

```

//run= g++ -fopenmp HPC_Practical_1.cpp -o HPC_Practical_1
//Windows:- HPC_Practical_1.exe
//Linux:- ./HPC_Practical_1

```

//Write a program to implement Parallel Bubble Sort and Merge sort using OpenMP.
Use existing algorithms and measure the performance of sequential and parallel algorithms.

```
#include <iostream>
#include <vector>
#include <omp.h>
using namespace std;
```

// Sequential Bubble Sort

```
void bubbleSortSeq(vector<int>& arr) {
    int n = arr.size();
    for (int i = 0; i < n-1; i++)
        for (int j = 0; j < n-i-1; j++)
            if (arr[j] > arr[j+1])
                swap(arr[j], arr[j+1]);
}
```

// Parallel Bubble Sort

```
void bubbleSortPar(vector<int>& arr) {
    int n = arr.size();
    for (int i = 0; i < n; i++) {
        #pragma omp parallel for
        for (int j = i % 2; j < n - 1; j += 2)
            if (arr[j] > arr[j + 1])
                swap(arr[j], arr[j + 1]);
    }
}
```

// Merge function

```
void merge(vector<int>& arr, int l, int m, int r) {
    vector<int> left(arr.begin() + l, arr.begin() + m + 1);
    vector<int> right(arr.begin() + m + 1, arr.begin() + r + 1);
    int i = 0, j = 0, k = l;
    while (i < left.size() && j < right.size())
        arr[k++] = (left[i] < right[j]) ? left[i++] : right[j++];
    while (i < left.size()) arr[k++] = left[i++];
    while (j < right.size()) arr[k++] = right[j++];
}
```

// Sequential Merge Sort

```
void mergeSortSeq(vector<int>& arr, int l, int r) {
    if (l < r) {
        int m = (l + r) / 2;
        mergeSortSeq(arr, l, m);
        mergeSortSeq(arr, m + 1, r);
        merge(arr, l, m, r);
    }
}
```

// Parallel Merge Sort

```
void mergeSortPar(vector<int>& arr, int l, int r) {
    if (l < r) {
        int m = (l + r) / 2;
```

```

        #pragma omp parallel sections
        {
            #pragma omp section
            mergeSortPar(arr, l, m);
            #pragma omp section
            mergeSortPar(arr, m + 1, r);
        }
        merge(arr, l, m, r);
    }
}

int main() {
    vector<int> data = {8, 5, 2, 9, 1, 4};
    vector<int> arr1 = data, arr2 = data;
    vector<int> arr3 = data, arr4 = data;

    bubbleSortSeq(arr1);
    bubbleSortPar(arr2);
    mergeSortSeq(arr3, 0, arr3.size() - 1);
    mergeSortPar(arr4, 0, arr4.size() - 1);

    cout << "Sorted (Seq Bubble): ";
    for (int x : arr1) cout << x << " ";
    cout << "\nSorted (Par Bubble): ";
    for (int x : arr2) cout << x << " ";
    cout << "\nSorted (Seq Merge): ";
    for (int x : arr3) cout << x << " ";
    cout << "\nSorted (Par Merge): ";
    for (int x : arr4) cout << x << " ";
}

```

//Compile Windows:- g++ -fopenmp HPC_Practical_2.cpp -o HPC_Practical_2
//Run:- HPC_Practical_2.exe

//Compile linux:- sudo apt update
//sudo apt install g++ libomp-dev
//g++ -fopenmp HPC_Practical_2.cpp -o HPC_Practical_2
//./HPC_Practical_2

bqcofxpye

April 18, 2025

0.1 Implement Min, Max, Sum and Average operations using Parallel Reduction.

```
[1]: import multiprocessing
import random

def parallel_reduction(operation, arr):
    with multiprocessing.Pool() as pool:
        if operation == "min":
            return min(pool.map(min, arr))
        elif operation == "max":
            return max(pool.map(max, arr))
        elif operation == "sum":
            return sum(pool.map(sum, arr))
        elif operation == "avg":
            return sum(pool.map(sum, arr)) / len(arr)

if __name__ == "__main__":
    arr = [random.randint(0, 10000) for _ in range(10000)]
    chunked_arr = [arr[i::multiprocessing.cpu_count()] for i in
↳range(multiprocessing.cpu_count())]

    print(f"Min: {parallel_reduction('min', chunked_arr)}")
    print(f"Max: {parallel_reduction('max', chunked_arr)}")
    print(f"Sum: {parallel_reduction('sum', chunked_arr)}")
    print(f"Average: {parallel_reduction('avg', chunked_arr)}")
```

```
Min: 0
Max: 10000
Sum: 49891056
Average: 6236382.0
```

```
[ ]:
```

5gj205fai

April 18, 2025

0.1 Write a cuda program for

1. Addition of two large vectors
2. Matrices Multiplication using CUDA C

```
[1]: import numpy as np
import time

n = 10**6
A = np.ones(n, dtype=np.float32)
B = np.full(n, 2.0, dtype=np.float32)

start = time.time()
C = A + B
end = time.time()

print(f"CPU vector add in {end - start:.6f} sec | C[0]: {C[0]}, C[-1]: {C[-1]}")
```

CPU vector add in 0.000999 sec | C[0]: 3.0, C[-1]: 3.0

```
[2]: import numpy as np
import time

# Matrix size
N = 1024

# Generate two random matrices
A = np.random.rand(N, N).astype(np.float32)
B = np.random.rand(N, N).astype(np.float32)

start = time.time()
C = np.matmul(A, B)
end = time.time()

print(f"CPU Matrix Multiplication Time: {end - start:.4f} seconds")
print(f"Result shape: {C.shape}")
```

CPU Matrix Multiplication Time: 0.0090 seconds
Result shape: (1024, 1024)